

Projet English



WORDPRESS



Sommaire


Context	3
Wordpress	4
Mysql	5
phpMyAdmin	6
Metric	7
DDOS	8
ISSUES	8
Deployment the file yaml	9
Glossary	10

Context

The final mission is developing a project that implements all of the concepts in this course by group of 2 people. We have a specifications to realise the project :

- Use kubernetes
- Containerize an application
- Provide a deployment manual with all necessary files
- Provide a report of the project

With all this information, we had to think about an application that would meet these expectations, so we decided to focus on Wordpress. We wanted to approach an application that could be used in a business context. We have set up a schedule to meet the deadlines by April 13th. Here is the schedule that we have set up:



To have complete management of Wordpress, we added metrics to the Kubernetes server, installed the PhpMyAdmin application to manage databases. Additionally, we added load management for Wordpress to connection overload.

Wordpress

WordPress is a free, open-source content management system used to create websites and blogs. It is one of the most popular platforms in the world and is known for its flexibility, ease of use, and large community of users and developers who contribute to its ongoing development and improvement. With thousands of customizable themes and plugins, WordPress can be used to build a wide range of websites, from personal blogs to complex e-commerce sites.

```
1 # generate secret for files
2 secretGenerator:
3   - name: mysql-pass
4     literals:
5       - password=Axellilian85!
6 # List des resources Deploy
7 resources:
8   - mysql-deployment.yaml
9   - wordpress-deployment.yaml
```

This block is to generate the secret for files and declaring the resources for the file. File kustomization.yaml

This block create the service of wordpress. This file is wordpress-deployment.yaml

```
# Declaring service
apiVersion: v1
kind: Service
metadata:
  name: wordpress
  labels:
    app: wordpress
spec:
  ports:
    - port: 80
  selector:
    app: wordpress
    tier: frontend
  type: LoadBalancer
```

```
---
# Asked the resources Volume storage
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: wp-pv-claim
  labels:
    app: wordpress
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 20Gi
```

This block asked the resource the volume storage.

This block creates the pod of WordPress.

```
# Create the pod of wordpress
apiVersion: apps/v1
kind: Deployment
metadata:
  name: wordpress
  labels:
    app: wordpress
spec:
  selector:
    matchLabels:
      app: wordpress
      tier: frontend
  strategy:
    type: Recreate
  template:
    metadata:
      labels:
        app: wordpress
        tier: frontend
    spec:
      containers:
        - image: wordpress:4.8-apache
          name: wordpress
          resources:
            requests:
              cpu: 50m
          env:
            - name: WORDPRESS_DB_HOST
              value: wordpress-mysql
            - name: WORDPRESS_DB_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: mysql-pass
                  key: password
          ports:
            - containerPort: 80
              name: wordpress
          volumeMounts:
            - name: wordpress-persistent-storage
              mountPath: /var/www/html
      volumes:
        - name: wordpress-persistent-storage
          persistentVolumeClaim:
            claimName: wp-pv-claim
```

```

---
# Create autoscaling for the application Wordpress
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: wordpress
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: wordpress
  minReplicas: 1
  maxReplicas: 5
  targetCPUUtilizationPercentage: 5

```

Autoscaling allows you to create weights according to the required performance. It is in the file wordpress-deployment.

MySQL

MySQL is a popular open-source relational database management system (RDBMS) that is used to store, manage, and retrieve data. It is commonly used in web applications to store user data, product information, and other types of data.

```

# Create the service wordpress-mysql
apiVersion: v1
kind: Service
metadata:
  name: wordpress-mysql
  labels:
    app: wordpress
spec:
  ports:
    - port: 3306
  selector:
    app: wordpress
    tier: mysql
  clusterIP: None

```

This block creates the service wordpress-mysql.

It is in the file: mysql-deployment.yaml

This block asked the resource volume storage.

```

---
# Asked the resources Volume storage
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mysql-pv-claim
  labels:
    app: wordpress
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 20Gi

```

```

# Create the pod of wordpress-mysql
apiVersion: apps/v1
kind: Deployment
metadata:
  name: wordpress-mysql
  labels:
    app: wordpress
spec:
  selector:
    matchLabels:
      app: wordpress
      tier: mysql
  strategy:
    type: Recreate
  template:
    metadata:
      labels:
        app: wordpress
        tier: mysql
    spec:
      containers:
        - image: mysql:5.6
          name: mysql
          env:
            - name: MYSQL_ROOT_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: mysql-pass
                  key: password
          ports:
            - containerPort: 3306
              name: mysql
          volumeMounts:
            - name: mysql-persistent-storage
              mountPath: /var/lib/mysql
          volumes:
            - name: mysql-persistent-storage
              persistentVolumeClaim:
                claimName: mysql-pv-claim

```

This block creates the pod of wordpress-mysql.

phpMyAdmin

PhpMyAdmin is an open-source web application used for managing MySQL databases. It provides a user-friendly graphical interface that allows users to manage their databases without having to write SQL code. We have added this application because I would access of the database.

```
1 # Create the pod of phpmyadmin
2 apiVersion: apps/v1
3 kind: Deployment
4 metadata:
5   name: phpmyadmin-deployment
6   labels:
7     app: phpmyadmin
8 spec:
9   replicas: 1
10  selector:
11    matchLabels:
12      app: phpmyadmin
13  template:
14    metadata:
15      labels:
16        app: phpmyadmin
17    spec:
18      containers:
19        - name: phpmyadmin
20          image: phpmyadmin/phpmyadmin
21          ports:
22            - containerPort: 80
23          env:
24            # The host for read the data base
25            - name: PMA_HOST
26              value: wordpress-mysql.default.svc.cluster.local
27            - name: PMA_PORT
28              value: "3306"
29            - name: MYSQL_ROOT_PASSWORD
30              valueFrom:
31                secretKeyRef:
32                  name: mysql-secrets
33                  key: ROOT_PASSWORD
```

This block creates the pod of phpMyAdmin.

This block creates the service of phpmyadmin.

```
1 # Create the service phpmyadmin
2 apiVersion: v1
3 kind: Service
4 metadata:
5   name: phpmyadmin-service
6 spec:
7   type: LoadBalancer
8   selector:
9     app: phpmyadmin
10  ports:
11    - protocol: TCP
12      port: 8080
13      targetPort: 80
```

```
1 # generate secret for files
2 apiVersion: v1
3 kind: Secret
4 metadata:
5   name: mysql-secrets
6 type: Opaque
7 data:
8   ROOT_PASSWORD: Axellilian85
```

This file creates the secret for the configuration of phpMyAdmin.

Metric

Metrics Server is a Kubernetes component that collects and stores performance metrics from the nodes and pods of a Kubernetes cluster. It is used to monitor resource usage in the cluster and to help administrators make informed decisions about the capacity and efficiency of the cluster.

```
apiVersion: v1
kind: ServiceAccount
metadata:
  labels:
    k8s-app: metrics-server
  name: metrics-server
  namespace: kube-system
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  labels:
    k8s-app: metrics-server
  rbac.authorization.k8s.io/aggregate-to-admin: "true"
  rbac.authorization.k8s.io/aggregate-to-edit: "true"
  rbac.authorization.k8s.io/aggregate-to-view: "true"
  name: system:aggregated-metrics-reader
rules:
- apiGroups:
  - metrics.k8s.io
  resources:
  - pods
  - nodes
  verbs:
  - get
  - list
  - watch
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  labels:
    k8s-app: metrics-server
  name: system:metrics-server
rules:
- apiGroups:
  - ""
  resources:
  - nodes/metrics
  verbs:
  - get
- apiGroups:
  - ""
  resources:
  - pods
  - nodes
  verbs:
  - get
  - list
  - watch
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    k8s-app: metrics-server
  name: metrics-server
  namespace: kube-system
spec:
  selector:
    matchLabels:
      k8s-app: metrics-server
  strategy:
    rollingUpdate:
      maxUnavailable: 0
  template:
    metadata:
      labels:
        k8s-app: metrics-server
    spec:
      containers:
      - args:
        - --cert-dir=/tmp
        - --secure-port=4443
        - --kubelet-preferred-address-types=InternalIP,ExternalIP,Hostname
        - --kubelet-use-node-status-port
        - --kubelet-insecure-tls
        - --metric-resolution=1s
        image: registry.k8s.io/metrics-server/metrics-server:v0.6.3
        imagePullPolicy: IfNotPresent
        livenessProbe:
          failureThreshold: 3
          httpGet:
            path: /livez
            port: https
            scheme: HTTPS
            periodSeconds: 10
          name: metrics-server
        ports:
        - containerPort: 4443
          name: https
          protocol: TCP
        readinessProbe:
          failureThreshold: 3
          httpGet:
            path: /readyz
            port: https
            scheme: HTTPS
            initialDelaySeconds: 20
            periodSeconds: 10
          resources:
            requests:
              cpu: 100m
              memory: 200Mi
        securityContext:
          allowPrivilegeEscalation: false
          readOnlyRootFilesystem: true
          runAsNonRoot: true
          runAsUser: 1000
        volumeMounts:
        - mountPath: /tmp
          name: tmp-dir
      nodeSelector:
        kubernetes.io/os: linux
      priorityClassName: system-cluster-critical
      serviceAccountName: metrics-server
      volumes:
      - emptyDir: {}
        name: tmp-dir
```

```
---
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  labels:
    k8s-app: metrics-server
  name: metrics-server-auth-reader
  namespace: kube-system
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: Role
  name: extension-apiserver-authentication-reader
subjects:
- kind: ServiceAccount
  name: metrics-server
  namespace: kube-system
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  labels:
    k8s-app: metrics-server
  name: metrics-server:system:auth-delegator
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: system:auth-delegator
subjects:
- kind: ServiceAccount
  name: metrics-server
  namespace: kube-system
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  labels:
    k8s-app: metrics-server
  name: system:metrics-server
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: system:metrics-server
subjects:
- kind: ServiceAccount
  name: metrics-server
```

```
---
apiVersion: v1
kind: Service
metadata:
  labels:
    k8s-app: metrics-server
  name: metrics-server
  namespace: kube-system
spec:
  ports:
  - name: https
    port: 443
    protocol: TCP
    targetPort: https
  selector:
    k8s-app: metrics-server
```

```
---
apiVersion: apiregistration.k8s.io/v1
kind: APIService
metadata:
  labels:
    k8s-app: metrics-server
  name: v1beta1.metrics.k8s.io
spec:
  group: metrics.k8s.io
  groupPriorityMinimum: 100
  insecureSkipTLSVerify: true
  service:
    name: metrics-server
    namespace: kube-system
  version: v1beta1
  versionPriority: 100
```

DDOS

It is to test the autoscaling of the WordPress. To confirm the solution of autoscaling.

```
# Create pod to DDOS my application wordpress
apiVersion: apps/v1
kind: Deployment
metadata:
  name: web-connections
spec:
  replicas: 20
  selector:
    matchLabels:
      app: web-connections
  template:
    metadata:
      labels:
        app: web-connections
    spec:
      containers:
        - name: web-connections
          image: nginx
          command: ["/bin/sh", "-c", "while true; do curl -s http://172.28.4.112 >/dev/null; sleep 1; done"]
```

ISSUES

I have issues like:

- To add the autoscaling, I have an error code to the deployment because we don't have the request resource so I add three line to declared the request resource:

```
containers:
  - image: wordpress:4.8-apache
    name: wordpress
    resources:
      requests:
        cpu: 50m
```

- To add the metric server, I have added one line: --kubelet-insecure-tls- Do not verify the CA of serving certificates presented by Kubelets.
- For phpMyAdmin, I have modified two lines because we don't acces in website so I change the port and the access

```
# Create the service phpmyadmin
apiVersion: v1
kind: Service
metadata:
  name: phpmyadmin-service
spec:
  type: LoadBalancer
  selector:
    app: phpmyadmin
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80
```

- The autoscaling don't work because the resource was good so I create a service DDOS to test the autoscaling

Deployment the file yaml

Download file in GitHub:

> 221-PERSO (F:) > Licence > A2SR > Communication Anglaise > Projet				
Nom	Modifié le	Type	Taille	
Phpmyadmin	05/04/2023 13:58	Dossier de fichiers		
wordpress	05/04/2023 09:07	Dossier de fichiers		
components.yaml	10/04/2023 13:11	Fichier source Yaml	5 Ko	
Password.txt	29/03/2023 10:03	Document texte	1 Ko	
web-connections.yaml	06/04/2023 19:52	Fichier source Yaml	1 Ko	

```
PS F:\Licence\A2SR\Communication Anglaise\Projet> kubectl apply -f .\components.yaml
```

```
PS F:\Licence\A2SR\Communication Anglaise\Projet\wordpress> kubectl apply -k ./
```

```
PS F:\Licence\A2SR\Communication Anglaise\Projet\Phpmyadmin> kubectl apply -f .\
```

```
PS F:\Licence\A2SR\Communication Anglaise\Projet> kubectl apply -f .\web-connections.yaml
```

Glossary

apiVersion: the version of the Kubernetes API used for this service.

kind: the type of Kubernetes resource.

metadata: additional metadata for the service, such as the name of the service.

spec: the specification

selector: the labels to use for selecting the pods to expose as a service.

ports: the list of ports to expose on the service.

name: the name of the port.

port: the port number to use for the service.

targetPort: the port number to use for the selected pods.

type: the type of service to create. If chosen LoadBalancer, which will create an external IP address for accessing in web server.

containers: the list of containers to run in the pod, which includes:

name: the name of the container.

image: the container image to use for running the server.

command: the command to run to start the server and deploy.

ports: the list of ports the container should expose, which includes:

name: the name of the port.

containerPort: the port number on which the container is listening for incoming connections.

env: the environment variables to use for the container.